Break This App: an Ecologist's Intro to Shiny Part 1

Lyndsie S. Wszola

4/21/2020

Hello and welcome to my first Shiny tutorial! Shiny is a framework for building interactive webpages using the R programming language. That means you can build powerful, customizable online tools to engage stakeholders in research communication using code and analyses you've already developed. In this tutorial, we'll learn how to build a Shiny app that shows temporal trends in Stellar sealion pup counts at a user-specified count site over time. This tutorial is meant to be an absolute basic intro to Shiny, so if you're already making Shiny apps it might be a bit too basic for you. If you're getting ready to build your first app, you're in the right place!

What do I need to make a Shiny app?

To get started, you'll need to have R and RStudio installed on your computer. The main coding prerequisite to building our sealion Shiny app is a basic understanding of the Tidyverse, a collection of R packages for data science. If you're new to the Tidyverse, I highly recommend the book R for Data Science, as well as the unofficial solutions manual. If you haven't used Tidyverse before, no worries! We'll go over each function we use in detail.

Ready to get started? Awesome! Let's open RStudio and make sure our required packages are installed if you don't have them already:

```r
install.packages("tidyverse")
install.packages("shiny")
```

Getting started

To start a new shiny app, we would click File>New File>Shiny Web App, which would create a new app directory for us. Since we already have an app to work with, we'll instead download or clone this repository if you haven't already ("Download or clone" is just github speak for make a copy of this whole big folder on your computer!). We'll follow along with the "app.R file" in the SeaLionApp folder. When you open app.R, you'll see some code at the top that sets up our workspace. We put these lines outside the main app code because that way they're run and loaded into the environment just once, which reduces memory needs and run time. The first lines load our required package libraries, shiny and tidyverse:

```r
library(shiny)
library(tidyverse)
```

In the next lines, we set up a custom ggplot theme, which we'll call "simpleTheme" so we can easily control the way our plot looks later:

```
simpleTheme <- theme_bw()+
    theme(panel.grid = element_blank(),
          axis.text = element_text(size = 20,
                      color = "black"),
          axis.title = element_text(size = 24,
                        color = "black"))
```

Now we're ready to load our data! The purpose of the app is to show trends in pup count data at different sites, so we need a list of sites. The following line loads an R data object containing a named list of the different count sites.

```
siteChoices <- read_rds("data/siteChoices.rds")
```

A note on directories:

Sometimes directories can be a little tricky. You'll notice that the lines above use the path "data/" because a published Shiny app needs all the data to be in a folder called "data," which it automatically looks for in the app directory. However,if you download and modify this app, sometimes your paths may end up a little different. When you're running the app locally on your computer, it may be easier to use the commented out lines, which will help the app find the right paths to run on your computer:

```
# appPath <- dirname(rstudioapi::getSourceEditorContext()$path)
# siteChoices <- read_rds(paste0(appPath, "/data/siteChoices.rds"))
```

In that case, you could comment out

```
siteChoices <- read_rds("data/siteChoices.rds")
```

A note on data preparation:

This app need two pieces of data: the main dataframe containing sites, dates, and mean pup counts, and a named list of sites for users to choose from. The focus of this tutorial is on getting an app off the ground, so I prepared the data ahead of time as .rds files. The .rds file format saves a single R object, like a list or dataframe, to file. This lets us restore the data object with traits like factor levels intact. I like .rds files because they let us do our data cleaning outside the scope of the app. This keeps our app code tidy and prevents us from wasting memory on processing data that will always need to be in the same format. BUT, you absolutely don't need to put your data in .rds files when you're making your own apps. You can use whatever data format (e.g., .csv's, .tif) you like! You can even type lists and other vectors right into the server or ui file. If you want to check out how I processed the data for this app, check out the dataProcessing.R file in this repo.

Basic app structure

Shiny apps are built in two pieces: the user interface and server. The user interface (ui), as the name suggests, creates all the elements that the user engages with. It contains all the static text and defines the appeerence of interactive widgets like dropdowns and sliders, as well as page layouts. We define the user interface using the `shinyUI` function:

```
ui <- shinyUI()
```

The second fundamental element, the server, does all the background work to run the app. The "magic" of a Shiny app is reactivity, the feedback between server and ui: the user manipulates the interactive elements in

the ui, then the server process user input and returns dynamic data and graphics elements. We define the server using the `shinyServer()` function:

```
server <- shinyServer(function(input, output, session){

})
```

Shiny assembles the server and ui into a cohesive unit using the `shinyApp()` function:

```
shinyApp(ui = ui, server = server)
```

Now that we have a handle on the basic structure of the app, let's dive into our sealion visualization! In each step, we will build on the previous step, until we build all the code we need to run the app.

UI

The first thing we need to decide for our new app's ui is how we want the text and visualization to be organized on the pages. Remember that Shiny apps make websites, so we have to plan out both the overall structure of the app's pages and how they relate, as well as the layout and organization of each page. The organization of the user interface code is consequently hierarchical, with each page's elements nested in a page layout, which is in turn nested in the whole app's layout.

We'll start out with an app layout that creates clickable tabs across the top, called a "navbarPage." I like navbarPage because it looks clean and professional, but there are lots of other options. We define this structure using the `navbarPage()`. We'll define the title of our navbarPage, which will show up in the top left, by typing it in quotes: `"Stellar Sealion Pup Count Trends."` So we've updated our empty shinyUI function from above to read:

```
ui <- shinyUI(navbarPage("Stellar Sealion Pup Count Trends"))
```

Now we need to choose the layout of our first page, which we initiate using the `tabPanel()` function and assign it a title that will appear in the navbar at the top of the page. We're going to use a `sidebarLayout`, which creates a shaded a side panel and a main panel. The side panel and main panel can both hold text, plots, maps, user input widgets, videos, and whatever else we need them to hold. The advantage of using a sidebarPanel() is that it draws the user's eye, so it's a great structure for communicating instructions and introductory material. Our sidepanel will hold a brief welcome for the user, made large using the `h2()` function, and a user input widget that prompts users to select a pup count site from a dropdown menu.

We define the dropdown menu using the `selectInput()` function, which allows users to select one site option. We give the drop-down widget a name, "site," which will pass to the server, and give it some human-friendly text to display to the reader: `"Choose a site to display the time series of mean pup counts at that site."`. We then specify the site choices to display to the reader. The choices come from the `siteChoices` object we loaded at the beginning. `siteChoices` is a named list, which means that it's a list of the levels of site, and the site names for those levels. The numeric indexing is important for helping R find the right site name from user input, and adding the site names creates a quick way for us to ensure that we can find both the correct level and its name with concise coding. For example, our `selectInput` statement sets the default selection to 1, the first index. The name of the first item in the list is "All," and it is a summary of all the sites. When the user selects an option from the `selectInput()` function, the selection is passed to server via `input` as the list index of the site chosen.

In addition, we'll add a main panel, using the `mainPanel()` function, to the right of our sidebarPanel to hold our dynamically updating graph of pup trends. Inside the mainPanel, we just have our sealion plot, which we print using the `plotOutput()` function and the name of the plot, "trendPlot" (more on that later). So altogether, our user interface now looks like:

```r
ui <- shinyUI(navbarPage("Stellar Sealion Pup Count trends", # page title
              tabPanel("Graphing pup counts", # Tab panel title
               sidebarLayout( # give us a nice data selection side bar!
                sidebarPanel(
                 h2("Welcome to the Stellar Sea Lion vizualization app! "),
                 # Some nice welcoming text for users
                 br(), br(),

                 selectInput("site", # Make a drop-down menu
                  "Choose a site to display the
                  time series of mean pup counts at that site.",
                 # data selection prompt for the users
                  choices = siteChoices,
                 # referencing the siteChoices list we loaded above
                  selected = 1)
                 # Make a default selection, the first item in the list, "All" sites
                   ),

                 mainPanel( # The big panel next to the side panel

                 plotOutput("trendPlot") # Print the plot we make below.

                 )))
                ))
```

The data we're using in our app comes from NOAA Fisheries, which means that we need to cite it properly.
We'll build a "Citations" page using a tabPanel as we did for the "Graphing pup counts" tab. The "Citations"
page will contain a wellPanel, a shaded offset panel, telling users who collected the data (NOAA fisheries)
and who made the app (me), and provide guidance on appropriate use. The wellPanel will also link users to
the data source using a hyperlink written with html tags and a string that describes what to show the user
the wellPanel. This ability to mix syntaxes and borrow elements from html and other languages is one of the
coolest things about shiny. With the addition of our citation page, we now have our completed user interface!

```r
ui <- shinyUI(navbarPage("Stellar Sealion Pup Count trends", # page title
              tabPanel("Graphing pup counts", # Tab panel title
               sidebarLayout( # give us a nice data selection side bar!
                sidebarPanel(
                 h2("Welcome to the Stellar Sea Lion vizualizationapp!"),
                 # Some nice welcoming text for users
                 br(), br(),

                 selectInput("site", # Make a drop-down menu
                  "Choose a site to display the time
                   series of mean pup counts at that site.",
                  # data selection prompt for the users
                  choices = siteChoices,
                  # referencing the siteChoices list we loaded above
                  selected = 1)
                  # Make a default selection, the first item in the list, "All" sites
                  ),

                 mainPanel( # The big panel next to the side panel
```

```
                 plotOutput("trendPlot") # Print the plot we make below.

             ))),
         tabPanel("Citations",
          # Must cite the data collectors! In a nice new tab panel!

         wellPanel("This app was created by Lyndsie S. Wszola
                       at the University of Nebraska for instructional
                       purposes only. The Stellar sealion pupcount
                       dataset was collected by The National Marine
                       Fisheries Service. More information on
                       the original dataset can be found",


                    tags$a(href="https://www.fisheries.noaa.gov/resource/data/counts-alaska
                    -steller-sea-lion-pups-conducted-rookeries-alaska-1961-2015", "here. "),

                   "The sealion pupcount data has been slightly
                    summarized and aggregated for this app. The
                    data visualizaions contained in the app are
                    therefore purely for educational purposes
                    and do not represent any
                    official finding or policy.")
                    # We must have a disclaimer so that people
                    #don't misuse/misinterpret anything here.

     )

     ))
```

Server

Now that we have our user interface set up, we need some cool data products for it to display! We'll start out by loading the sealion pup count data as we did for the named list of sites above (the note about file paths applies here too).

```
server <- shinyServer(function(input, output, session) {

         # seaLions <- read_rds(paste0(appPath, "/data/seaLions.rdata"))

         seaLions <- read_rds("data/seaLions.rds")

   })
```

Now we need to build the links between the user interface and the data wrangling and plotting functions we're about to describe. If you've ever built functions that depend on other functions, the structure of a Shiny server is very similar. If you haven't, no worries! We'll walk through it step by step.

Reactivity

Reactivity is the ability of Shiny apps to receive user inputs and perform operations that are then returned to the user as updating visuals and values. This helps our Shiny app update only what's needed at each step and keep its processing time reasonable. Even though we only receive the user's input once, each time we use the information we receive from the user, it will need to be a `reactive({})` value because all the

values "downstream" of the user's input depend on the value of that input. It's generally good practice to keep our reactive values simple, so that when (not if) our code breaks, we will have an easier time identifying the problem.

Our first reactive value, `siteSelection`, will extract the name of the site specified by the user in the drop-down menu. This is why we went to all that trouble with a named list in a .rds file! If you recall the structure of our siteChoices list, it was a list of numeric values describing the levels of the pup count site variable. The list also has names, corresponding to the names of each site. We can access all the names of the list via `names(siteChoices)`. The user's site choice is passed as a numeric value to server via `input$site`. So, to access the name of the selected site, we extract from our vector of site names the name whose index corresponds to the user's selected site (Try saying that 5 times fast, or better yet thinking through it 1 time slowly). The last important detail in the siteSelection reactive value is wrapping input$site in an `as.numeric()`, just to make sure it gets interpreted as a numeric value.

```r
siteSelection <- reactive({

    site <- names(siteChoices)[as.numeric(input$site)]
    # The user's selection from the drop-down

})
```

Once the user's site selection has been passed into the reactive context, we need to filter our data to include only observations from the site the user selected. We will accomplish our data filtering using our `siteSelection()` reactive value from above (notice how it's written like a function) and the `filter()` function from the dplyr package. We start our dplyr statement by specifing the dataframe we want to use, `seaLions`. We then add a pipe symbol (%>%), which basically tells R "take everything that has happened up to this point in this code block, and continue it down to the next line." Once the user's site selection has been passed into the reactive context, we need to filter our data to include only observations from the site the user selected. We will accomplish our data filtering using our `siteSelection()` reactive value from above (notice how it's written like a function) and the `filter()` function from the dplyr package. We start our dplyr statement by specifing the dataframe we want to use, `seaLions`. We then add a pipe symbol (%>%), which tells R "take everything that has happened up to this point in this code block, and continue it down to the next line." We then use the `filter` function to keep only those rows in our dataframe whose "site" variable matches exactly the "site" reactive value defined above.

Notice that we use `dplyr::filter()` with the `::`, which indicates that we want to use the `filter()` function from the dplyr package. I did this because it's a good habit to specify packages in large coding projects you might share with other people, but also because we might want to add some spatial components in a future tutorial. Several spatial packages also have a `filter()` function, so specifying the package now saves us some future headaches.

```r
data <- reactive({ # Subsetting the data based on user input

    site <- siteSelection()
    # referencing site-selection from above

    siteDat <- seaLions %>%
                # Keep only data from the user-specified site.
                dplyr::filter(sitename == site)

})
```

Plotting

Now that we have our data reactively updating to match user specifications, it's time to make our graph! We bring our user-specified filtered data into the scope of the plot using the `data()` reactive value and assign it to an object called `siteDat`, which will operate in the scope of the plot. We will create a line plot with points on top of the lines using the `renderPlot({})` function, which we will designate as an output object by calling it `output$trendPlot`. If you refer back to the UI code above, you will notice that the mainPanel on the "Graphing pup counts" page displays this graph by calling `plotOutput(trendPlot)`.

We will build the pup count trend plot using the `ggplot()` function from the ggplot2 package, which we loaded as part of the tidyverse along with dplyr. ggplot graphics are built in layers, so each line of code will visually appear on top of the one before it. That means that each time we add a new graphics argument, it will inherit data and aesthetics (the x and y specification) unless we choose to overwrite it.

We start our pupcount trend plot with a line graph using the `geom_line()` function that creates a continuous line whose x coordinates are the years, specified `as.numeric()` just in case, and whose y coordinates are the mean pups counted in each year. We will then add points on top of the line using the `geom_point()` function, with the same x and y coordinates. The default point size is just a little smaller than I want it to be, so we'll make it bigger by specifying size = 3. Next, we specify the x and y limits of the plot. The `coord_cartesian()` function allows us to specify x and or y limits. We'll set the y axis to adjust with the data. By specifying htat `ylim = c(0, max(siteDat$meanPups + 5))`, we tell ggplot to limit the y axis to a minimum of 0 and a maximum of 5 pups above our maximum number of mean observed pups. Next, we set the limits of our x axis. We'll keep the years constant so the user can see the different time spans of data collection. We specify the limits and breaks, where the axis ticks should go, `using the scale_x_continuous()` function. Finally, we add some axis labels using the `labs()` function and add the theme we set up at the beginning to make the graph layout nice and simple.

```r
output$trendPlot <- renderPlot({

siteDat <- data() # Source our data from above

ggplot(siteDat) + # Make a beautiful (simple!) plot
    geom_line(aes(x = as.numeric(year),
                  y = meanPups)) +
    geom_point(aes(x = as.numeric(year),
                   y = meanPups), size = 3) +
    coord_cartesian(ylim = c(0, max(siteDat$meanPups + 5))) +
    scale_x_continuous(limits = c(1960, 2015),
                       breaks = seq(1960, 2020, 10)) +
    labs(x = "Year", y = "Mean pups per observation") +
    simpleTheme

        })
```

At this point, we should have a fully functional Shiny app that shows sealion pup count data at a site of the user's choosing. We can and will add more later, but it's a good start. We can run our new app using the "Run App" button that should have appeared in the upper right hand corner of your RStudio window. It will launch the app in a browser window.

Other Shiny resources

The Shiny app gallery

Mastering Shiny book

Why break this app?

I called this tutorial "Break this app" in the hopes that you would do just that: take the example app and break it. Then fix it, add your own elements to it, and adapt the basic structure for your project. The best way to learn Shiny is to simply change one or a few things at a time and observe the effects. Just pick something and change it. You can start simple: text color, size, or content, graph style, and then work up to things like changing page layouts and adding output. If it works, great! If it doesn't, you still learned something. Good luck and happy coding!